

Linux/Unix Overview and Access Control

Victor Francisco Fonte
University of Minho
vff@di.uminho.pt

2015

Linux 101

- Multi-program, multi-task, multi-user
- Kernel, libraries, system calls, processes
- Super-user, regular user
- Shell, system and user-level command-line programs
- Configs and service management
- Optional GUI and graphical apps

Linux 101 (2)

- Program vs. process
- Virtual memory with address space isolation
- File system
- Inter-process communication

Processes

- Running program
- Isolated address space
- Executed under an user and group id
 - actual or virtual users
 - real or effective user and group ids
- May spawn other processes

Process: address space

- Traditional linear model
- Code (lower), data, and stack (higher)
- Data: initialised vs. non-initialised (e.g. heap)
- Access to the null (zero) address is illegal

Process: file descriptors

- File descriptor represents a kernel-level file-like object
- Processes inherited file descriptors: standard input, output, and error
- Processes inherit all opened file descriptors from their fathers

Process: creation and execution

- Processes are created by cloning
 - address space is copied (lazy)
 - opened file descriptors are inherited
- Programs are executed
 - address space is replaced by the contents of an executable file (ABI)
 - opened file descriptors are preserved

Inter-process Communication

- Files
- Traditional IPC:
 - shared memory
 - message queues
 - semaphores
- Pipes (named or un-named), sockets (local or network domains), ...

Unix File and Directory Permissions and Modes

- POSIX-compliant systems
- POSIX 1:2008 a.k.a. Single Unix Specification (SUS) V4

Users, Groups and Passwords

- User ID (UID), group ID (GID)
- Primary and secondary groups
- User DB: /etc/passwd (and /etc/shadow)
- Group DB: /etc/group

Access Control to Resources

- Homogeneity, simplicity
- Files, directories, and other resources
- Ownership and permissions

Ownership

- Assigned user owner
- Assigned group owner

Classes of Users

1. User owner
2. Member of the owning group
3. Others (everyone else)

Checked by this order.

Permissions

- Read, write, execute
- For each class of user
- Permissions set independently
- Not all combinations make sense
- Files and directories are similar
- But slightly different semantics

Permissions

- umask: user-default restriction on permissions
- 12 bit permissions kept in resource i-node
- Permissions influence syscalls (not commands)
- Only super-user can modify ownership
- Only owner (or super-user) can modify perms

Permissions set on Files

- Read: access content
- Write: modify content
- Execute: execute content

Directories

- Directories are special files
- Map names to i-nodes
- Access to names protected by 'read'
- Access to i-nodes protected by 'execute'

Permissions set on Dirs

- Read: view/list names
 - Write: add/delete/rename names
 - Execute: chdir, access i-nodes (stat)
-
- Execute a.k.a search permission
 - Execute permissions needed along a path
 - Permissions not inherited from parent dir

Permissions and other Attributes

- perms, u-owner, g-owner, c/m-date, name, ...
- perms: type, perms, set uid/gid, sticky, ...
- ex: -rwxr--r--, drwxr-xr-x
- ex: drwsrwsrwt, drwSrwsrwT
- ex: -rwxrw-r--+

Permissions: Experimentation

- `chmod`, `chown`, `chgrp`
- `read()`, `write()`, `exec()`, `unlink()`, `stat()`
- `strace`

Permissions: Experimentation

- `mkdir adir && touch adir/afile`
- `ls adir; ls -l adir; ls -l adir/afile`
- `chmod -x adir`
- `ls adir; ls -l adir; ls -l adir/afile`

Real vs. Effective User & Group IDs

- Default: (EUID, EGID) = (RUID, RGID)
- If SUID, SGID are set on resources
 - EUID = resource user owner ID
 - EGID = resource group owner ID

SUID & SGID Attrs

- Useful but dangerous
- Violates the minimal privileges principle
- Lookout for root SUID
- In particular, root SUID + 'write' permission
- Executable/Not-executable: 's'/'S'

SUID & SGID on Files

- 'S' on user owner id: no meaning
- 'S' on group owner id: mandatory locking
- 's' has no effect on scripts... why?

SUID & SGID on DLLs

- Shared objects: `.so`, `.so.number`
- System-wide defaults: `/etc...`
- Override: `LD_LIBRARY_PATH`
- Dangerous for SUID, SGID... why?
- If SUID or SGID then `LD_LIBRARY_PATH` is ignored

Sticky Bit on Files

- Sticky a.k.a. Text Bit
- Code kept on swap or memory
- Deprecated in favor of virtual memory
- Executable/Non-executable: 't'/'T'
- No effect on non-executables

SUID & SGID on Dirs

- SUID no effect on directories
- On Linux & Solaris SGID: group owner copied to entries (no inheritance!)
- SGID semantics is not in SUS V4

Sticky Bit on Dirs

- If set, only owner of a resource, the owner of the directory (or the super-user) can move/rename/delete the resource
- Works in conjunction with 'write' perm

Access Control Lists

- Permission set to specific users and groups
- Complements the UGO/RWX mechanism
- Default ACL for directories can be inherited, but usually are just copied to entries
- POSIX proposal withdrawn
- Solution to the per directory umask problem

ACLs: Experimentation

- Utilities: setfac, getfac
- Effective vs real permission mask
- Denoted by a '+' after perms

Additional/Extended Attributes

- Additional attributes:
 - NTFS, ext-family, ...
 - ext2: lsattr, chattr
- Extended attributes:
 - name, value pairs
 - getfattr, setfattr
- A ' ' or ' ' after perms unless ACL is set

Rootly Powers

- Some operations ignore permissions
- Check who makes the request: EUID = 0?
- Eg. system shutdown, port binding, ...
- Solutions: capabilities mechanism
- Program start as root but give up capabilities they don't need
- Minimize privileges to exploit

Chroot Jail

- changes the apparent root directory
- affects running process and its children
- operation restricted to the super-user
- usage: testing and development, dependency control and compatibility, recovery, privilege separation

Chroot best practices

- change working directory into the jail before chroot
- change real/effective user id to a non-root
- keep as little as possible inside the jail
- have root own as many jailed read-only files as possible
- limit all permissions of files and directories
- create a permissions-setting script
- chroot from inside the daemon itself (avoid wrapping)

Chroot best practices

- preload dynamically loaded objects
- avoid using the jailed `/etc/passwd` file
- close file descriptors aggressively before chrooting
- link config files from the outside
- update environment variables to reflect the new root

Minimal chroot

- `close(unused file descriptors);`
- `chdir("jail");`
- `chroot(".");`
- `setuid(non-root-uid);`

Chroot

- related system calls: `chroot()`
- related commands: `chroot`
- precursor of jail, Solaris Containers, Linux Containers, Docker, Linux userspaces (kernel 3.8+)

Further reading

- http://en.wikipedia.org/wiki/Unix_security
- http://www.tldp.org/LDP/intro-linux/html/sect_03_04.html
- <http://www.cse.psu.edu/~trjl/cse497b-s07/slides/cse497b-lecture-18-unixsecurity.pdf>